Table of model training choices

	Optimizing Accuracy (Xception)	**Optimizing Latency (MobileNetV2)**
Data Transformation/Augmentation	Random rotation, zoom, width shift, height shift, horizontal flip	Random rotation, zoom, width shift, height shift, horizontal flip
Base Model (include name, size, top-1 accuracy, CPU inference time))	Name: Xception Size: 88 MB Top-1 Accuracy: 79.0% CPU Inference Time: 109.4 ms	Name: MobileNetV2 Size: 14 MB Top-1 Accuracy: 71.3% CPU Inference Time: 25.9 ms
Number of epochs, Optimizer, and learning rate used to train classification head	20 epochs, ADAM @ 0.01 (early stopping stopped epochs at 4)	5 epochs, ADAM @ 0.01
number of layers un-frozen	5	5
Number of epochs, Optimizer, and learning rate used to further fine-tune the model	30 epochs, Adam @ 0.0001	10 epochs, ADAM @ 0.0001
final accuracy on evaluation set (test set)	89.75%	86.47%

Performance of models when deployed as a single pod

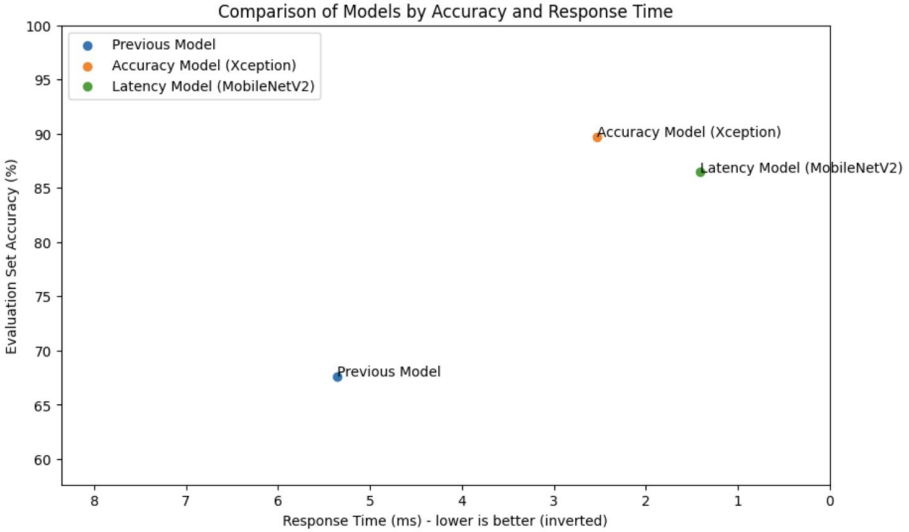

Legend
Previous Model
Accuracy Model (Xception)
Latency Model (MobileNetV2)

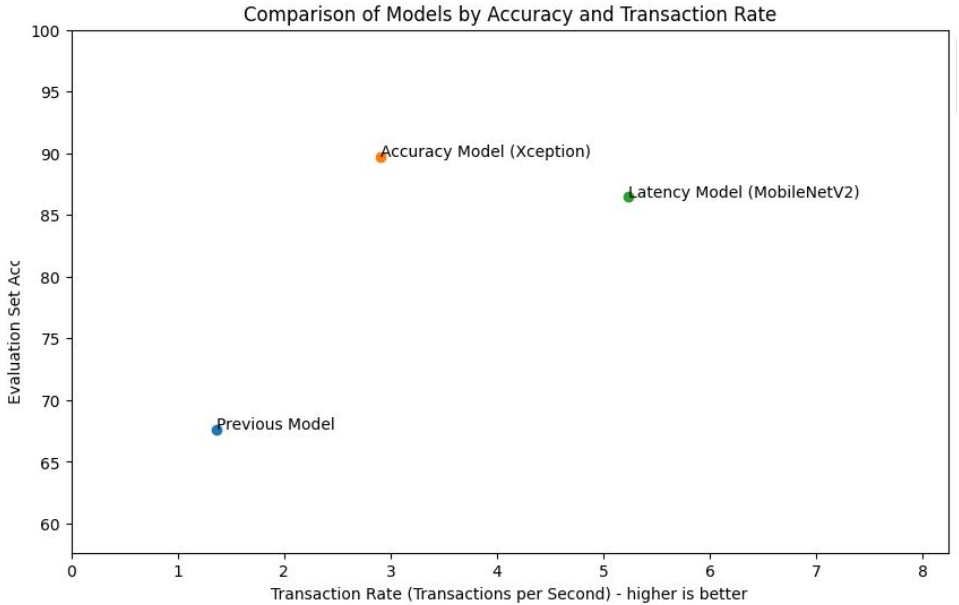

Performance of models when deployed as a "max-size" deployment

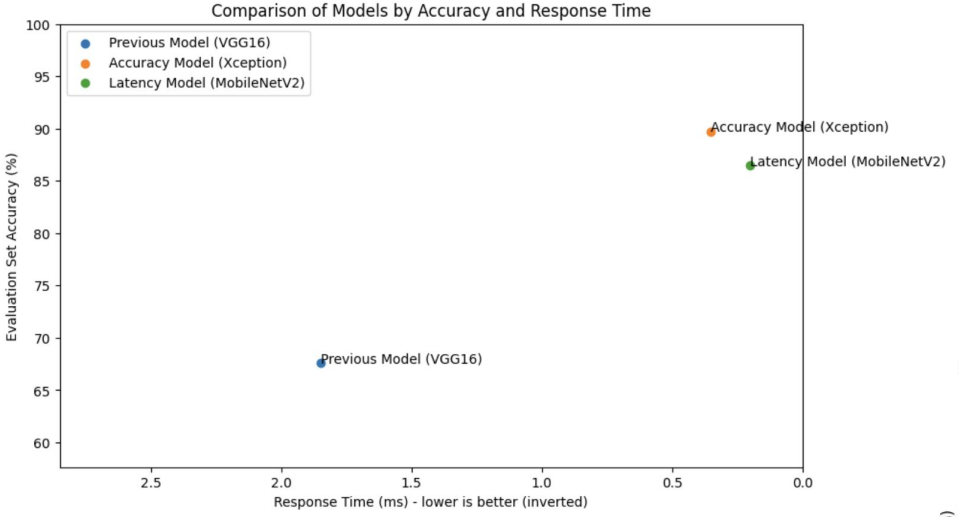

Legend
Previous Model
Accuracy Model (Xception)
Latency Model (MobileNetV2)

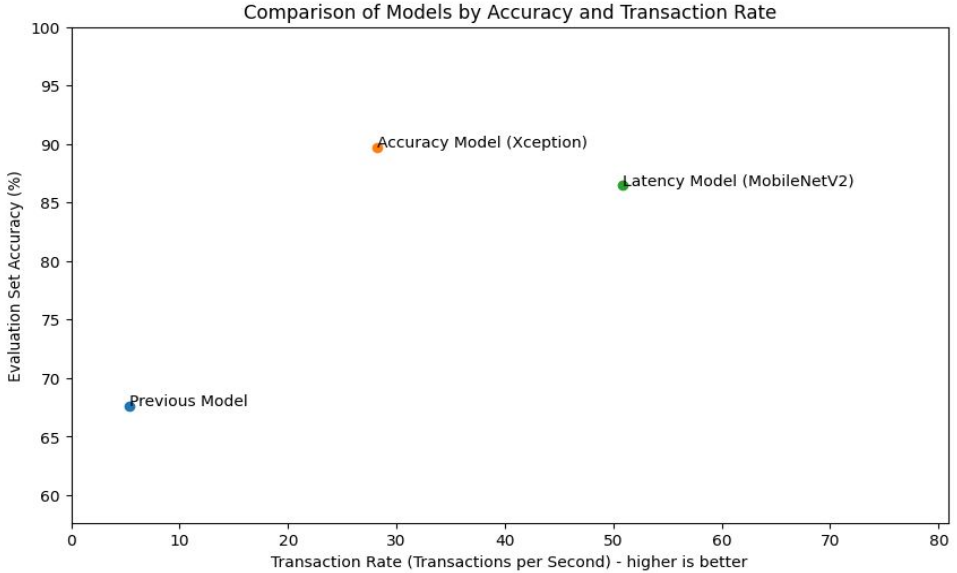

Table showing replicas and resource request configurations for "max-size" deployment

Criteria	Previous	Accuracy (Xception)	Latency (MobileNet)
Number of replicas	20	12	9
CPU resource requests	0.1	0.2	1
Memory resource requests	0.5Gi	1Gi	15 Mi
CPU resource limits	2	2	2
Memory resource limits	4Gi	4Gi	4Gi

Table showing horizontal scaling configurations

	Accuracy (Xception)	Latency (MobileNetV2)
minReplicas	3	4
maxReplicas	10	15
targetCPUUtilizationPercentage	75%	40%
CPU resource requests	3.2 cores	0.85
Memory resource requests	6Gi	1Gi
CPU resource limits	4 cores	4 cores
Memory resource limits	8Gi	2Gi

Visualization of deployment for "accuracy" model over time

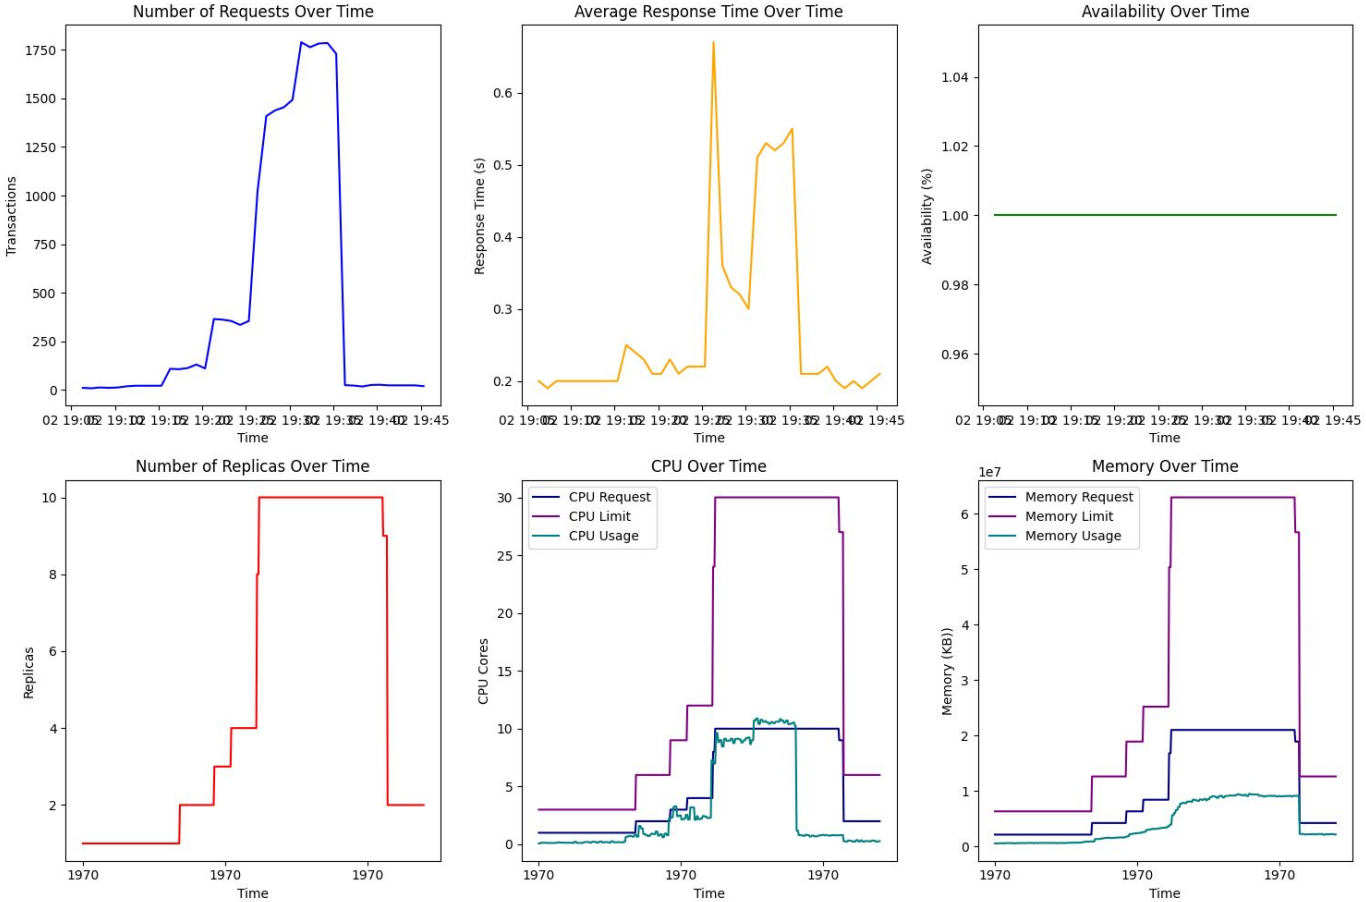

Visualization of deployment for "latency" model over time

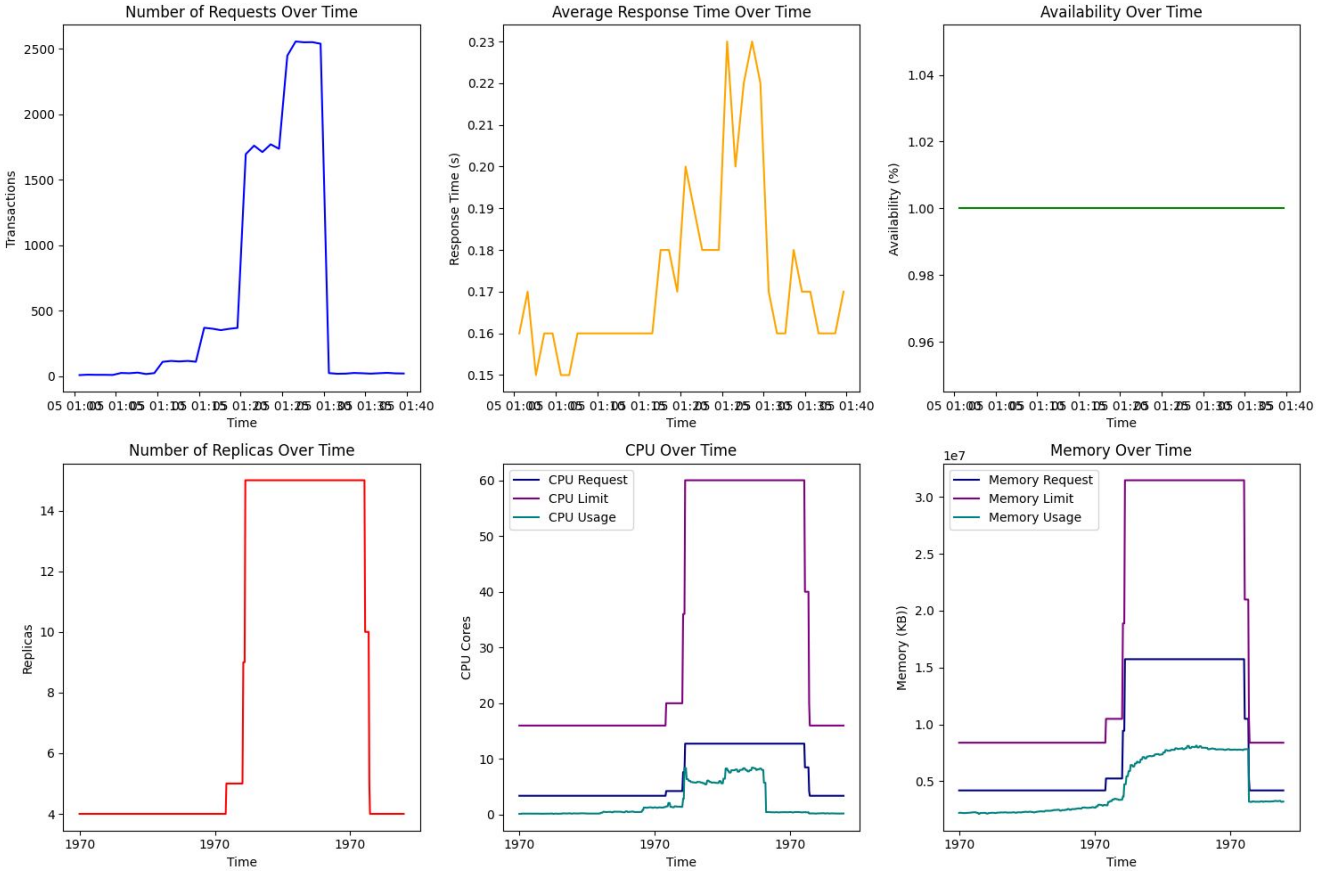

Summarize your contributions

All in all,

- The **previous model** has low accuracy (67.64%) and high inference time (1.28 seconds) . When deployed, It has high response time, low availability, low transaction rate. **There is much scope for improvement.**
- We implement **Xception** Model to focus on accuracy. Which has good accuracy & good enough model size. We implement horizontal scaling and set configurations such that the system is **highly available and scales during high requests**. We see that it's accuracy is **89.75% and inference time is (1.19 seconds)**.
- We implement **MobileNetV2** Model to focus on latency. Which has better accuracy than previous model & very less inference time per step. We implement horizontal scaling and **allocate more resources such it reduces inference time** and scales during high requests. We see that it's accuracy is **86.74% and inference time is (0.05 seconds).**